


## 9.5 Bucket and Sample Sort

A popular serial algorithm for sorting an array of  $n$  elements whose values are uniformly distributed over an interval  $[a, b]$  is the **bucket sort** algorithm. In this algorithm, the interval  $[a, b]$  is divided into  $m$  equal-sized subintervals referred to as **buckets**, and each element is placed in the appropriate bucket. Since the  $n$  elements are uniformly distributed over the interval  $[a, b]$ , the number of elements in each bucket is roughly  $n/m$ . The algorithm then sorts the elements in each bucket, yielding a sorted sequence. The run time of this algorithm is  $\Theta(n \log(n/m))$ . For  $m = \Theta(n)$ , it exhibits linear run time,  $\Theta(n)$ . Note that the reason that bucket sort can achieve such a low complexity is because it assumes that the  $n$  elements to be sorted are uniformly distributed over an interval  $[a, b]$ .

Parallelizing bucket sort is straightforward. Let  $n$  be the number of elements to be sorted and  $p$  be the number of processes. Initially, each process is assigned a block of  $n/p$  elements, and the number of buckets is selected to be  $m = p$ . The parallel formulation of bucket sort consists of three steps. In the first step, each process partitions its block of  $n/p$  elements into  $p$  sub-blocks, one for each of the  $p$  buckets. This is possible because each process knows the interval  $[a, b]$  and thus the interval for each bucket. In the second step, each process sends sub-blocks to the appropriate processes. After this step, each process has only the elements belonging to the bucket assigned to it. In the third step, each process sorts its bucket internally by using an optimal sequential sorting algorithm.

Unfortunately, the assumption that the input elements are uniformly distributed over an interval  $[a, b]$  is not realistic. In most cases, the actual input may not have such a distribution or its distribution may be unknown. Thus, using bucket sort may result in buckets that have a significantly different number of elements, thereby degrading performance. In such situations an algorithm called **sample sort** will yield significantly better performance. The idea behind sample sort is simple. A sample of size  $s$  is selected from the  $n$ -element sequence, and the range of the buckets is determined by sorting the sample and choosing  $m - 1$  elements from the result. These elements (called **splitters**) divide the sample into  $m$  equal-sized buckets. After defining the buckets, the algorithm proceeds in the same way as bucket sort. The performance of sample sort depends on the sample size  $s$  and the way it is selected from the  $n$ -element sequence.

Consider a splitter selection scheme that guarantees that the number of elements ending up in each bucket is roughly the same for all buckets. Let  $n$  be the number of elements to be sorted and  $m$  be the number of buckets. The scheme works as follows. It divides the  $n$  elements into  $m$  blocks of size  $n/m$  each, and sorts each block by using quicksort. From each sorted block it chooses  $m - 1$  evenly spaced elements. The  $m(m - 1)$  elements selected from all the blocks represent the sample used to determine the buckets. This



**Figure 9.20** An example of the execution of sample sort on an array with 24 elements on three processes.

scheme guarantees that the number of elements ending up in each bucket is less than  $2n/m$  (Problem 9.28).

How can we parallelize the splitter selection scheme? Let  $p$  be the number of processes. As in bucket sort, set  $m = p$ ; thus, at the end of the algorithm, each process contains only the elements belonging to a single bucket. Each process is assigned a block of  $n/p$  elements, which it sorts sequentially. It then chooses  $p - 1$  evenly spaced elements from the sorted block. Each process sends its  $p - 1$  sample elements to one process – say  $P_0$ . Process  $P_0$  then sequentially sorts the  $p(p - 1)$  sample elements and selects the  $p - 1$  splitters. Finally, process  $P_0$  broadcasts the  $p - 1$  splitters to all the other processes. Now the algorithm proceeds in a manner identical to that of bucket sort. This algorithm is illustrated in Figure 9.20.

**Analysis** We now analyze the complexity of sample sort on a message-passing computer with  $p$  processes and  $O(p)$  bisection bandwidth.

The internal sort of  $n/p$  elements requires time  $\Theta((n/p) \log(n/p))$ , and the selection of  $p - 1$  sample elements requires time  $\Theta(p)$ . Sending  $p - 1$  elements to process  $P_0$  is similar to a gather operation (Section 4.4); the time required is  $\Theta(p^2)$ . The time to internally sort the  $p(p - 1)$  sample elements at  $P_0$  is  $\Theta(p^2 \log p)$ , and the time to select  $p - 1$  splitters is  $\Theta(p)$ . The  $p - 1$  splitters are sent to all the other processes by using one-to-all broadcast (Section 4.1), which requires time  $\Theta(p \log p)$ . Each process can *insert* these  $p - 1$  splitters in its local sorted block of size  $n/p$  by performing  $p - 1$  binary searches. Each process thus partitions its block into  $p$  sub-blocks, one for each bucket. The time required for

this partitioning is  $\Theta(p \log(n/p))$ . Each process then sends sub-blocks to the appropriate processes (that is, buckets). The communication time for this step is difficult to compute precisely, as it depends on the size of the sub-blocks to be communicated. These sub-blocks can vary arbitrarily between 0 and  $n/p$ . Thus, the upper bound on the communication time is  $O(n) + O(p \log p)$ .

If we assume that the elements stored in each process are uniformly distributed, then each sub-block has roughly  $\Theta(n/p^2)$  elements. In this case, the parallel run time is

$$T_P = \overbrace{\Theta\left(\frac{n}{p} \log \frac{n}{p}\right)}^{\text{local sort}} + \overbrace{\Theta\left(p^2 \log p\right)}^{\text{sort sample}} + \overbrace{\Theta\left(p \log \frac{n}{p}\right)}^{\text{block partition}} + \overbrace{\Theta(n/p) + O(p \log p)}^{\text{communication}}. \quad (9.9)$$

In this case, the isoefficiency function is  $\Theta(p^3 \log p)$ . If bitonic sort is used to sort the  $p(p-1)$  sample elements, then the time for sorting the sample would be  $\Theta(p \log p)$ , and the isoefficiency will be reduced to  $\Theta(p^2 \log p)$  (Problem 9.30).